

agent = model + harness; the harness is everything below

THE MODEL: reasoning only (Claude session, Mode B) — no memory between calls, no ground truth

1 run the loop

burst-walkthrough
protocol: RUN →
PRESENT → GATE
→ LITERATURE →
DISTILL

dispatcher: roster,
gates, order at task
intake

prototype chains:
fit, merge, products,
temporal, assembly

queue manager +
typed workflows (A17)

the repository layer

2 execute tools

spectral engine:
24-model menu ×
every usable bin
(campaign run),
validity + rail columns

Bayesian blocks (3ML
native) + significance
merge

Stage-1 approval GUI
+ ingest: detectors,
background, source —
decision human or AI,
stamped

figures, temporal
chain, SED family,
per-burst report

3 manage context

27 guide files; 10 step
skills indexed by the
ledger; lessons with
IDs; memory files +
notes

skill-reader: the
binding checklist at
every step open

standing contracts in
the PI's quoted words
(S-, R-items)

isolation: each verifier
and reader runs in its
own fresh context

4 persist state

state board derived
from DISK evidence,
never from memory

identity stamps,
receipts, sidecars,
commit pin

object schemas (10) +
validation test; action
trace with model id

first-class actions:
intent → preflight →
commit → receipt

5 enforce boundaries

hooks: no-ship
(sha), dispatch plan,
raw-data guard

producer ≠ verifier ≠
approver; fresh-context
verifiers

fit-table auditor,
tie rule, count
coordinates, edge
stamps

loop caps, spend
ledger, post-tool audit,
pre-commit, sandbox

PI gate: every step

identity-stamped into the records above; never fabricated

hosted coding product: runtime loop, tool calls, compaction, permissions

code: deterministic,
tested

model judgement: a
role in fresh context

documents, records:
read or written, not
run

human

(dashed: PROPOSED)

solid = on disk at the census
(2026-09-02); each box maps
to rows of Table T1; hatched
frames = the two harness layers;
grey headers = the five jobs